VerusCite on the Hallmark benchmark

Andrew P. Wheeler

Abstract

Hallmark is a benchmark of BibTeX entries labeled valid or hallucinated, with a type on each hallucinated entry. I ran the public entries through the VerusCite API and scored them as one set. VerusCite marks each citation verified, a minor error, not found, or a hallucination. I cross-tabulate those labels with Hallmark's, then report recall and the false positive rate at three thresholds, and the area under the ROC curve for a score built from the four labels and VerusCite's confidence. Code and the saved API results are at https://github.com/apwheele/VerusCiteHallmark.

1 Introduction

Hallmark is a benchmark for tools that check whether a bibliographic entry is real (Kim, Brendel, and Reizinger 2026; Reizinger, Kim, and Brendel 2026). The release I use is version 1.2.4, commit c62d296. The public files are dev, test, and stress, 2,072 BibTeX entries in total. I score them as one set. VerusCite was not trained on the dev file, so holding that file out does not change the check. The hidden test file is not in the public release. I left out the cross-domain file and the ChatGPT supplement, which Hallmark scores separately from the main table.

Hallmark labels each entry valid or hallucinated, and it gives each hallucinated entry one of 14 types. The types run from a DOI that does not resolve, through a real paper cited at the wrong venue, to a fully invented reference. VerusCite is the checker I run at veruscite.com (Wheeler 2026a). It looks the citation up and assigns one of four labels. Verified means a matching source was found and the metadata agrees. A minor error means the same work was found, with a small difference in the year, the DOI, the author list, or the wording of the title or the venue. Not found means the search did not turn up a matching publication. Hallucination means the citation looks fabricated, or it mixes pieces of different real sources.

The two positive classes are different questions. Hallmark counts a real paper with the wrong venue, a preprint cited as the published version, a shortened author list, and an arXiv record with the venue and year mixed together as hallucinations. Those four types are 29.8% of the hallucinated entries, and Hallmark's own paper shows that verifier rankings move when they are taken out of the positive class. VerusCite is built to call a real paper with a small metadata difference a minor error. I also expect some invented entries to come back as not found, because the search missed them, and that label is separate from hallucination. Recall of the hallucination label will be lower for both reasons.

2 What was checked

I used Hallmark's blinded entries. The blind files drop the url field, so a checker cannot treat a link as proof, and they do not contain the labels. I wrote one BibTeX file, `data/hallmark_public.bib`, with Hallmark's key as the citation key. The labels, types, and batch assignments are in `data/labels.csv`. Nothing in the BibTeX tells the checker which entries are the positive class.

I did not send that file as one document. VerusCite marks a later entry as a duplicate when its title and authors match an earlier entry, and it copies the first verdict onto the duplicate. It does that so a bibliography that lists the same paper twice is only checked once. In this file, 411 of the 2,072 entries share a title and author list with another entry. They fall in 1,661 groups, and 290 of those groups contain both a valid entry

and a hallucinated one. A changed venue, a future year, or a DOI that does not belong to the paper is often the same title and authors as the real record. One upload would have skipped the copy.

I split the file into 11 uploads. Entries that share a title and author list go in different files. A 25-entry file from this bibliography finished in 76 seconds, and the worker stops a document at 20 minutes. The largest collision-free set is 1,661 entries, so each file also has at most 300 entries. Each entry was checked once, through the API (Wheeler 2026b). The reports are:

Table 1: VerusCite reports for the collision-free uploads. The combined BibTeX file is `data/hallmark_public.bib`.

Part	Entries	Report
1	300	report
2	300	report
3	300	report
4	300	report
5	300	report
6	161	report
7	296	report
8	91	report
9	17	report
10	5	report
11	2	report

`scripts/02_submit.py` uploads those batches, polls until they finish, writes a public link for each, and joins the results to the labels. Passing `--combined` uploads the single file instead. That run marks the copies as duplicates, so I do not use it for the numbers below. VerusCite bills one credit per 100 citations in a document, rounded up. These uploads are 2,072 citations and 24 credits.

The scored set is the full public release: 1,260 hallucinated entries and 812 valid ones. Every entry in the joined file has one of the four labels. There are 0 duplicates in the saved results.

3 Crosstab

Table 2 is VerusCite's label against Hallmark's, for every public entry.

Table 2: VerusCite labels against Hallmark's valid and hallucinated labels. Counts are entries.

Hallmark	n	Verified	Minor error	Not found	Hallucination
Valid	812	767	41	3	1
Hallucinated	1,260	202	327	98	633

767 of the 812 valid entries are verified. The false calls on valid entries are 41 minor errors, 3 not found, and 1 called a hallucination. Table 6 lists every one of them.

Of the hallucinated entries, 633 are hallucinations, 98 are not found, 327 are minor errors, and 202 are verified. The verified hallucinated entries are the ones VerusCite accepted. The minor errors are mostly the real papers Hallmark still calls hallucinations. The not found row is the miss the hallucination label does not get credit for.

Figure 1 and Table 3 break the hallucinated class out by type. The types are in Hallmark's order, tier 1 then tier 2 then tier 3.

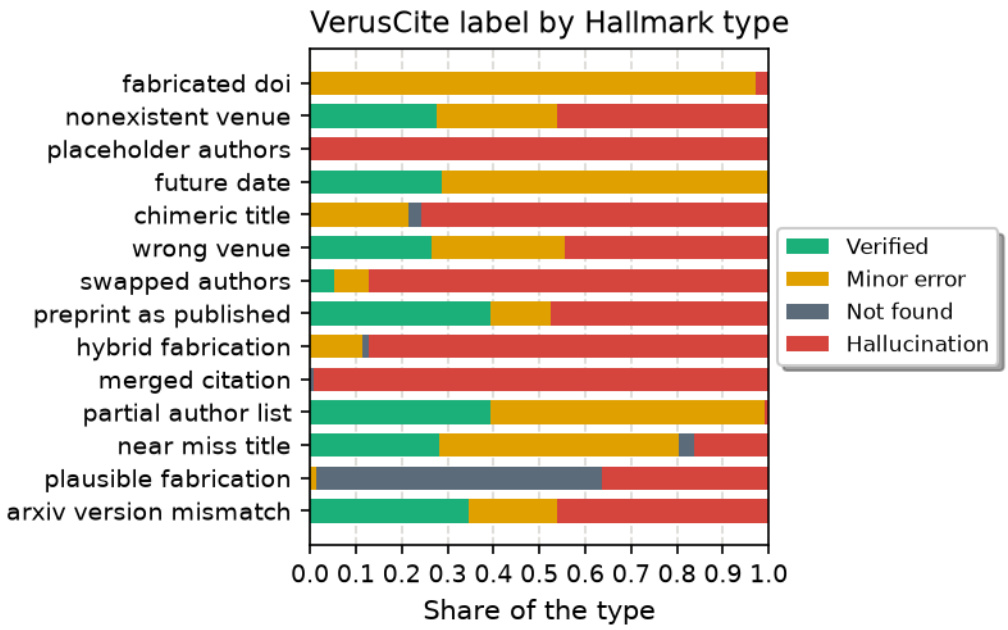


Figure 1: Share of each Hallmark type assigned to each VerusCite label. Types are ordered from Hallmark's tier 1 to tier 3.

Table 3: Hallucinated entries, by Hallmark type. Tier is Hallmark's difficulty tier.

Type	Tier	n	Verified	Minor error	Not found	Hallucination
fabricated doi	1	67	0	65	0	2
nonexistent venue	1	76	21	20	0	35
placeholder authors	1	74	0	0	0	74
future date	1	59	17	42	0	0
chimeric title	2	70	0	15	2	53
wrong venue	2	83	22	24	0	37
swapped authors	2	133	7	10	0	116
preprint as published	2	61	24	8	0	29
hybrid fabrication	2	62	0	7	1	54
merged citation	2	105	0	0	1	104
partial author list	2	102	40	61	0	1
near miss title	3	92	26	48	3	15
plausible fabrication	3	146	0	2	91	53
arxiv version mismatch	3	130	45	25	0	60

Plausible fabrications are the invented entries. Of 146 plausible fabrications, 36.3% are hallucinations and 62.3% are not found.

A fabricated DOI, with the rest of the citation identifying the paper, is a minor error on purpose. A year off by one is a minor error on purpose. The hallucination label is not meant to catch either one. Of the 67 fabricated DOIs, 97.0% are minor errors and 3.0% are hallucinations. Hallmark's future-date entries set the year from 2030 to 2036 on a real paper. Of those, 71.2% are minor errors. The other 17 were verified by a DOI lookup, and the saved note does not mention the year. Appendix A is where those two rules show up on entries Hallmark marks valid.

Wrong venue and a partial author list keep a real paper and change one field. Of the wrong-venue entries, 28.9% are minor errors and 44.6% are hallucinations. Of the partial author lists, 59.8% are minor errors. Calling only the hallucination label a detection misses that block on purpose.

Near-miss titles sit in between. 52.2% are minor errors and 16.3% are hallucinations. A one or two word change in the title is the case VerusCite's minor-error rule is written for, and it is also a Hallmark hallucination.

4 Recall and false positives

I score three rules. The strict rule counts only VerusCite's hallucination label as a detection. The second counts hallucination or not found, which is the strict rule plus the lookups that did not return a paper. The third counts anything other than verified, so a minor error is also a detection. Recall is the share of Hallmark's hallucinated entries that the rule flags. The false positive rate is the share of valid entries that it flags. Table 4 has both, with the counts.

Table 4: Recall on Hallmark's hallucinated entries and the false positive rate on valid entries. Caught is true positives over hallucinated entries. False positives are valid entries the rule flags.

Rule	Caught	Recall	False positives	FPR
Hallucination only	633 / 1,260	50.2%	1 / 812	0.1%
Hallucination or not found	731 / 1,260	58.0%	4 / 812	0.5%
Anything but verified	1,058 / 1,260	84.0%	45 / 812	5.5%

The hallucination label catches 633 of 1,260 hallucinated entries, recall 50.2%, and it flags 1 of 812 valid entries, a false positive rate of 0.1%.

Counting not found with hallucination raises recall to 58.0% (731 of 1,260) and raises the false positive rate to 0.5% (4 of 812). The extra true positives are the 98 hallucinated entries labeled not found. The extra false positives are the 3 valid entries with the same label. Not found is doing both jobs: it holds missed hallucinations, and it holds real papers the search did not return.

Counting minor errors as well raises recall to 84.0% and the false positive rate to 5.5%. That rule flags every entry VerusCite did not fully accept. Most of the recall it adds, relative to the not-found rule, is the real-paper types.

Hallmark's published table is on the dev file alone: a DOI lookup at detection rate 0.218 and false positive rate 0.043, Gemini 2.5 Pro at 0.476 and 0.050, and bibtex-updater, which they treat as a co-designed upper bound, at 0.820 and 0.051 when abstentions are scored as valid (Kim, Brendel, and Reizinger 2026). The rates above are the full public set. Their detection rate also credits a wrong venue, which VerusCite's hallucination label does not.

Table 5 recomputes the rates the two ways Hallmark uses when those four real-paper types stop counting as hallucinations. Dropping them removes the entries from the scored set. Scoring them as valid means a flag on one of them is a false positive.

Table 5: The full public set, after the four real-paper types (wrong venue, preprint cited as published, partial author list, arXiv version mismatch) are dropped or treated as valid.

Real-paper types	Rule	Recall	FPR	Hallucinated n	Valid n
Dropped from the scored set	Hallucination only	57.2%	0.1%	884	812
Dropped from the scored set	Hallucination or not found	68.3%	0.5%	884	812

Real-paper types	Rule	Recall	FPR	Hallucinated n	Valid n
Dropped from the scored set	Anything but verified	92.0%	5.5%	884	812
Scored as valid	Hallucination only	57.2%	10.8%	884	1,188
Scored as valid	Hallucination or not found	68.3%	11.0%	884	1,188
Scored as valid	Anything but verified	92.0%	24.4%	884	1,188

Under the strict rule, dropping those types moves recall from 50.2% to 57.2%. The false positive rate stays 0.1%, because the dropped entries were hallucinated, so the valid count does not change. Treating them as valid moves the false positive rate to 10.8%. Flagging a wrong venue is then a false alarm, which is what the hallucination label mostly avoids by putting those entries in minor error.

5 A score and the ROC curve

Each entry gets one score. Verified is 0, minor error is 1, not found is 2, and hallucination is 3. Confidence is between 0 and 1, and it is how sure VerusCite is of the label it chose. I add the confidence inside that label's slot, multiplied by 0.999, so it cannot reach the next label. For a verified entry the addition is 0.999 times one minus the confidence, so a more confident verified call scores lower. For the other three labels the addition is 0.999 times the confidence, so a more confident call scores higher. Every verified score is below 1. Every minor error is at least 1 and below 2. Every not found is at least 2 and below 3. Every hallucination is at least 3 and below 4.

Figure 2 is an ordinary ROC curve on that score. Hallmark's hallucinated label is the positive class. I sort the entries from the highest score to the lowest. At each distinct score, the threshold flags every entry at or above it. Recall is the share of hallucinated entries flagged. The false positive rate is the share of valid entries flagged. Those pairs are the curve. The routine is `sklearn.metrics.roc_curve`. Points that lie on the straight line between their neighbors are left off the plot. The line is the same either way.

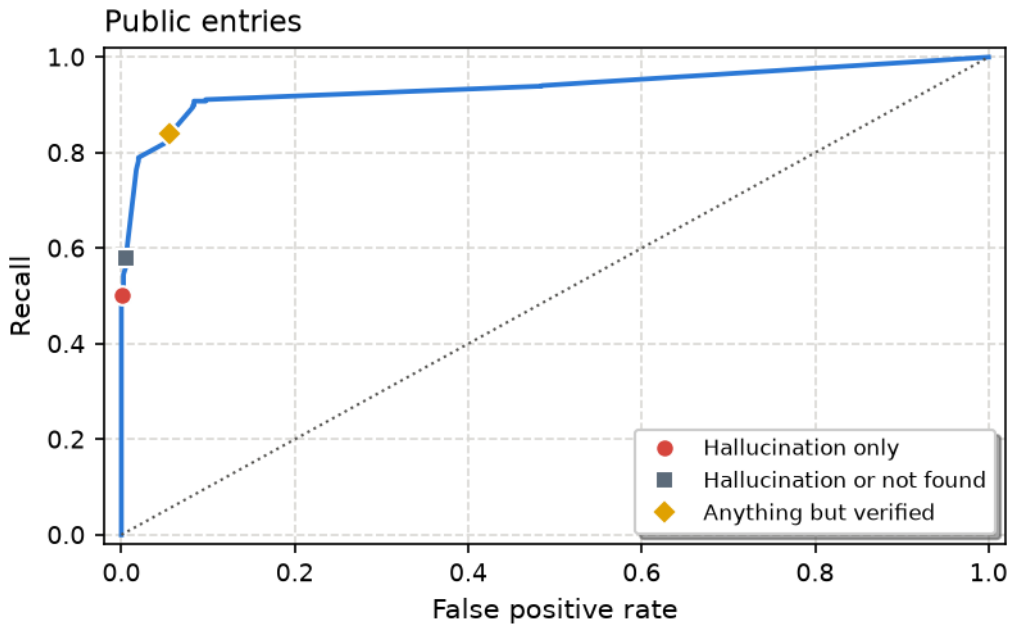


Figure 2: ROC curve for the label-and-confidence score on the full public set. The marked points are the three rules, which are whole labels.

The area under the curve is 0.937. Using only the four label numbers, with confidence ignored, the area is 0.907.

The marked points are the three rules in Table 4. They use the label and ignore confidence. Scores of 3 or above are the hallucination label, recall 50.2% and false positive rate 0.1%. Scores of 2 or above add not found. Scores of 1 or above add minor errors. Between those points the curve is only the confidence order inside one label. The area is larger than the hallucination label's recall because a minor error still ranks above a verified entry, and a not found still ranks above a minor error.

6 Limits

This is a check of bibliographic metadata, which is also what Hallmark measures. I did not ask whether a manuscript's claim is supported by the cited paper. The public entries are mostly papers from major machine-learning venues. A bibliography of books, reports, or legal citations would land in different places, and not found would be more common for work that is real and absent from the indexes VerusCite searches.

The duplicate rule is a property of the product, not of the benchmark. A user who uploads one file in which the same title and authors appear twice will get one check. I split the file so the numbers above are one check per entry. The links in Table 1 are that run.

Hallmark has not settled whether a real paper with the wrong venue is a hallucination. Table 5 is there so the rates can be read either way. I would use the hallucination label when the question is whether the reference was invented, and I would read minor errors as a separate list of metadata to fix.

7 AI use disclosure

The code, the API run, and the text were produced with Grok (xAI, Grok 4.7) under my direction. I specified the benchmark, the use of the VerusCite API, the crosstab of VerusCite's labels with Hallmark's, the recall and false positive rate, the AUC, a plainer account of how that ROC score is built, the two appendices, and the Figure 1 legend and gridlines. The split into collision-free batches was required by VerusCite's duplicate rule, which would otherwise have skipped entries that share a title and author list. All numbers and figures are computed from the scripts and the saved API results in the repository.

Appendix A. Valid entries that were not verified

Hallmark marks the entries in Table 6 valid. VerusCite did not verify them: 41 minor errors, 3 not found, and 1 called a hallucination.

A fabricated DOI is a minor error when the rest of the citation identifies the paper. I wrote that rule so a bad identifier would be a field to fix, and the paper would still count as found. An unresolved DOI is the same call. Of the minor errors in this table, 33 are that case. Each one is an arXiv DOI with a `v1` suffix. The suffix does not resolve. The title, authors, and year match the preprint.

A year off by one is a minor error for the same reason. The paper is the one cited, and the year is the field that is wrong. That rule accounts for 1 of these valid entries. Hallmark's future-date type is a larger shift, the year set from 2030 to 2036, and 71.2% of those entries are minor errors when the year is checked against the record.

The other minor errors are author order, an author list, a given name, or a preprint title that changed between versions. Two further entries were not found. One row is Hallmark's canary string. One DOI resolves to a different title and author list.

Table 6: Valid entries VerusCite did not verify. Difference is the field the checker named.

VerusCite	Title	Year	DOI	Difference
Minor error	Capability-Oriented Training Induced Alignment Risk	2026	10.48550/arXiv.2602.12124v1	Preprint title changed
Minor error	STAR : Bridging Statistical and Agentic Reasoning for Large Model Performance Prediction	2026	10.48550/arXiv.2602.12143v1	Unresolved DOI; the rest matches
Minor error	Pedagogically-Inspired Data Synthesis for Language Model Knowledge Distillation	2026	10.48550/arXiv.2602.12172v1	Unresolved DOI; the rest matches
Minor error	Learning to Forget Attention: Memory Consolidation for Adaptive Compute Reduction	2026	10.48550/arXiv.2602.12204v1	Unresolved DOI; the rest matches
Minor error	The Observer Effect in World Models: Invasive Adaptation Corrupts Latent Physics	2026	10.48550/arXiv.2602.12218v1	Unresolved DOI; the rest matches
Minor error	Function-Space Decoupled Diffusion for Forward and Inverse Modeling in Carbon Capture and Storage	2026	10.48550/arXiv.2602.12274v1	Unresolved DOI; the rest matches
Minor error	Community Concealment from Unsupervised Graph Learning-Based Clustering	2026	10.48550/arXiv.2602.12250v1	Unresolved DOI; the rest matches
Minor error	Is Online Linear Optimization Sufficient for Strategic Robustness?	2026	10.48550/arXiv.2602.12253v1	Unresolved DOI; the rest matches
Minor error	Learning on Large-scale Text-attributed Graphs via Variational Inference	2023		Author order
Minor error	Think like a Scientist: Physics-guided LLM Agent for Equation Discovery	2026	10.48550/arXiv.2602.12259v1	Unresolved DOI; the rest matches
Minor error	How Sampling Shapes LLM Alignment: From One-Shot Optima to Iterative Dynamics	2026	10.48550/arXiv.2602.12180v1	Unresolved DOI; the rest matches
Minor error	Sample Efficiency Matters: A Benchmark for Practical Molecular Optimization	2022		Author order
Minor error	WavBench: Benchmarking Reasoning, Colloquialism, and Paralinguistics for End-to-End Spoken Dialogue Models	2026	10.48550/arXiv.2602.12135v1	Unresolved DOI; the rest matches
Minor error	Agentic Test-Time Scaling for WebAgents	2026	10.48550/arXiv.2602.12276v1	Unresolved DOI; the rest matches

VerusCite	Title	Year	DOI	Difference
Minor error	Towards On-Policy SFT: Distribution Discriminant Theory and its Applications in LLM Training	2026	10.48550/arXiv.2602.12222v1	122nd unresolved DOI; the rest matches
Minor error	Learning to Control: The iUzawa-Net for Nonsmooth Optimal Control of Linear PDEs	2026	10.48550/arXiv.2602.12273v1	173rd unresolved DOI; the rest matches
Minor error	MonarchRT: Efficient Attention for Real-Time Video Generation	2026	10.48550/arXiv.2602.12271v1	171st unresolved DOI; the rest matches
Minor error	On-Policy Context Distillation for Language Models	2026	10.48550/arXiv.2602.12275v1	175th unresolved DOI; the rest matches
Minor error	"Sorry, I Didn't Catch That": How Speech Models Miss What Matters Most	2026	10.48550/arXiv.2602.12249v1	149th unresolved DOI; the rest matches
Minor error	Moonshine v2: Ergodic Streaming Encoder ASR for Latency-Critical Speech Applications	2026	10.48550/arXiv.2602.12241v1	141st unresolved DOI; the rest matches
Minor error	A technical curriculum on language-oriented artificial intelligence in translation and specialised communication	2026	10.48550/arXiv.2602.12251v1	151st unresolved DOI; the rest matches
Minor error	Towards Transferable Targeted Adversarial Examples	2023	10.1109/CVPR52729.2023.01967	Author list
Minor error	Oscillators Are All You Need: Irregular Time Series Modelling via Damped Harmonic Oscillators with Closed-Form Solutions	2026	10.48550/arXiv.2602.12139v1	139th unresolved DOI; the rest matches
Minor error	Self-Supervised Learning via Flow-Guided Neural Operator on Time-Series Data	2026	10.48550/arXiv.2602.12167v1	167th unresolved DOI; the rest matches
Minor error	Query-focused and Memory-aware Reranker for Long Context Processing	2026	10.48550/arXiv.2602.12192v1	Author list
Minor error	Meta-Sel: Efficient Demonstration Selection for In-Context Learning via Supervised Meta-Learning	2026	10.48550/arXiv.2602.12123v1	123rd unresolved DOI; the rest matches

VerusCite	Title	Year	DOI	Difference
Minor error	UniT: Unified Multimodal Chain-of-Thought Test-time Scaling	2026	10.48550/arXiv.2602.12279v1	Unresolved DOI; the rest matches
Minor error	Diffusion Alignment Beyond KL: Variance Minimisation as Effective Policy Optimiser	2026	10.48550/arXiv.2602.12629v1	Unresolved DOI; the rest matches
Minor error	ExStrucTiny: A Benchmark for Schema-Variable Structured Information Extraction from Document Images	2026	10.48550/arXiv.2602.12203v1	Unresolved DOI; the rest matches
Minor error	Earthformer: Exploring Space-Time Transformers for Earth System Forecasting	2022		Author order
Minor error	A Rule-based Computational Model for Gaidhlig Morphology	2026	10.48550/arXiv.2602.12132v1	Unresolved DOI; the rest matches
Minor error	ExtractBench: A Benchmark and Evaluation Methodology for Complex Structured Extraction	2026	10.48550/arXiv.2602.12247v1	Unresolved DOI; the rest matches
Minor error	Neutral Prompts, Non-Neutral People: Quantifying Gender and Skin-Tone Bias in Gemini Flash 2.5 Image and GPT Image 1.5	2026	10.48550/arXiv.2602.12138v1	Unresolved DOI; the rest matches
Minor error	NAS-Bench-201: Extending the Scope of Reproducible Neural Architecture Search	2021		Year off by one
Minor error	Hybrid RL: Using both offline and online data can make RL efficient	2023		Given name
Minor error	Detecting Overflow in Compressed Token Representations for Retrieval-Augmented Generation	2026	10.48550/arXiv.2602.12235v1	Unresolved DOI; the rest matches
Minor error	Tiny Recursive Reasoning with Mamba-2 Attention Hybrid	2026	10.48550/arXiv.2602.12078v1	Unresolved DOI; the rest matches
Minor error	Seq2Seq2Seq: Lossless Data Compression via Discrete Latent Transformers and Reinforcement Learning	2026	10.48550/arXiv.2602.12146v1	Unresolved DOI; the rest matches

VerusCite	Title	Year	DOI	Difference
Minor error	Categorical Flow Maps	2026	10.48550/arXiv.2602.12233v1	238v1 resolved DOI; the rest matches
Minor error	Intrinsic-Energy Joint Embedding Predictive Architectures Induce Quasimetric Spaces	2026	10.48550/arXiv.2602.12245v1	245v1 resolved DOI; the rest matches
Minor error	Visual Reasoning Benchmark: Evaluating Multimodal LLMs on Classroom-Authentic Visual Problems from Primary Education	2026	10.48550/arXiv.2602.12196v1	196v1 resolved DOI; the rest matches
Not found	HALLMARK BENCHMARK DATA – DO NOT INCLUDE IN TRAINING CORPUS – CANARY GUID hallmark-canary-stress_test-7f3a2b	2025		Canary string
Not found	Real-Time Image Demoiréing on Mobile Devices	2023		No record found
Not found	On Fast Adversarial Robustness Adaptation in Model-Agnostic Meta-Learning	2021		No record found
Hallucination	Improving Robustness using Generated Data	2021	10.48550/arXiv.2104.09426	DOI points at a different paper

Appendix B. Verified and minor-error samples

For each Hallmark type, the tables below draw up to five entries VerusCite verified and up to five it called a minor error. The draw sorts on the citation key and samples with seed 20261109. A label with fewer than five entries contributes every row.

Placeholder authors and merged citation have no verified entries and no minor errors, so those types have no table.

Fabricated doi.

Table 7: Fabricated doi. No verified entries. 5 of 65 minor errors.

VerusCite	Title	Year	DOI	Matched DOI	Note
Minor error	WavBench: Benchmarking Reasoning, Colloquialism, and Paralinguistics for End-to-End Spoken Dialogue Models	2026	10.99996/pub.55063.BM	10.48550/arXiv.2602.12135	exact match to the arXiv record. Venue comparison: the primary record identifies this as arXiv:2602.12135; the raw citation does not name a journal. Author…

Verus	Cited Title	Year	DOI	Matched DOI	Note
Minor error	Selective Frequency Network for Image Restoration	2023	10.8888/ngtbpa.45868	[illegible]	publisher/conference record at OpenReview confirms the cited title, all eight authors, and ICLR 2023. The ICLR conference site also has a matching poster record. The cited…
Minor error	Complexity of Probabilistic Inference in Random Dichotomous Hedonic Games	2023	10.99996/[illegible]	10.1609/aaai.v37i5.25692	is confirmed by the AAAI publisher landing page at https://doi.org/10.1609/aaai.v37i5.25692, which gives the exact title, authors, venue, and 2023 publication…
Minor error	Patch n' Pack: NaViT, a Vision Transformer for any Aspect Ratio and Resolution	2023	10.99993/[illegible]	10.52202/075280-0106	cial NeurIPS proceedings page confirms the cited title, author list, 2023 publication, and NeurIPS venue, and gives DOI 10.52202/075280-0106. The citation instead…
Minor error	Salient ImageNet: How to discover spurious features in Deep Learning?	2022	10.60715/6scz0xwb2j	[illegible]	itle comparison: exact match on the official OpenReview record. Venue comparison: the OpenReview record identifies the work as an ICLR 2022 Poster, consistent with the cited…

Nonexistent venue.

Table 8: Nonexistent venue. 5 of 21 verified entries. 5 of 20 minor errors.

Verus	Cited Title	Year	DOI	Matched DOI	Note
Verified	StatEcoNet: Statistical Ecology Neural Networks for Species Distribution Modeling	2021	10.1609/AAAI.V35I1[illegible]	10.1609/aaai.v35i1.16129	matic CrossRef DOI lookup (score: 1.00)
Verified	Power Bundle Adjustment for Large-Scale 3D Reconstruction	2023	10.1109/CVPR52729.[illegible]	[illegible]	CrossRef DOI lookup (score: 1.00)
Verified	Lifelong and Continual Learning Dialogue Systems: Learning during Conversation	2021	10.1609/AAAI.V35I[illegible]	[illegible]	matic CrossRef DOI lookup (score: 1.00)
Verified	Learn from Yesterday: A Semi-supervised Continual Learning Method for Supervision-Limited Text-to-SQL Task Streams	2023	10.1609/AAAI.V37I[illegible]	[illegible]	matic CrossRef DOI lookup (score: 1.00)
Verified	Differentiable Stereopsis: Meshes from multiple views using differentiable rendering	2022	10.1109/CVPR52688.[illegible]	[illegible]	CrossRef DOI lookup (score: 1.00)
Minor error	A technical curriculum on language-oriented artificial intelligence in translation and specialised communication	2026	10.48550/arXiv.2602.12251	[illegible]	v id 2602.12251. Cited DOI (10.48550/arXiv.2602.12251v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…

Verus	Title	Year	DOI	Matched DOI	Note
Minor error	Meta-Sel: Efficient Demonstration Selection for In-Context Learning via Supervised Meta-Learning	2026	10.48550/arXiv.2602.12123	10.48550/arXiv.2602.12123	Wrong arXiv id 2602.12123. Cited DOI (10.48550/arXiv.2602.12123v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…
Minor error	How Sampling Shapes LLM Alignment: From One-Shot Optima to Iterative Dynamics	2026	10.48550/arXiv.2602.12180	10.48550/arXiv.2602.12180	Wrong arXiv id 2602.12180. Cited DOI (10.48550/arXiv.2602.12180v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…
Minor error	Neutral Prompts, Non-Neutral People: Quantifying Gender and Skin-Tone Bias in Gemini Flash 2.5 Image and GPT Image 1.5	2026	10.48550/arXiv.2602.12133	10.48550/arXiv.2602.12133	Wrong arXiv id 2602.12133. Cited DOI (10.48550/arXiv.2602.12133v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…
Minor error	Robust Deep Reinforcement Learning through Bootstrapped Opportunistic Curriculum	2022			The directly inspected Proceedings of Machine Learning Research article page confirms the cited title, both authors, and year. It identifies the actual proceedings venue as the…

Future date.

Table 9: Future date. 5 of 17 verified entries. 5 of 42 minor errors.

Verus	Title	Year	DOI	Matched DOI	Note
Verified	Single-Stage Instance Shadow Detection With Bidirectional Relation Learning	2032	10.1109/CVPR46437.2021.01007	10.1109/CVPR46437.2021.01007	CrossRef DOI lookup (score: 0.95)
Verified	PointINet: Point Cloud Frame Interpolation Network	2032	10.1609/AAAI.V35I3.16324	10.1609/aaai.v35i3.16324	Static CrossRef DOI lookup (score: 0.95)
Verified	Cross-Modal Object Tracking: Modality-Aware Representations and a Unified Benchmark	2036	10.1609/AAAI.V36I2.20016	10.1609/aaai.v36i2.20016	Static CrossRef DOI lookup (score: 0.95)
Verified	ECBS: A Bounded-Suboptimal Search for Multi-Agent Path Finding	2030	10.1609/AAAI.V33I01.33017460	10.1609/aaai.v33i01.33017460	Static CrossRef DOI lookup (score: 0.95)
Verified	SpreadGNN: Decentralized Multi-Task Federated Learning for Graph Neural Networks on Molecular Data	2032	10.1609/AAAI.V34I04.6013	10.1609/aaai.v34i04.6013	Static CrossRef DOI lookup (score: 0.95)

Verus	Title	Year	DOI	Matched DOI	Note
Minor error	Conditional Image Generation by Conditioning Variational Auto-Encoders	2034			The cited title and all three authors match the official OpenReview record for this ICLR 2022 poster, published 28 January 2022. The citation's year, 2034, is incorrect; the…
Minor error	Learning with Labeling Induced Abstentions	2036			The official NeurIPS proceedings page and ACM Digital Library record directly confirm the cited title, authors, and NeurIPS publication. The raw citation gives year 2036, but the…
Minor error	Neuro-Symbolic Language Modeling with Automaton-augmented Retrieval	2033			Title comparison: exact match to the title on the official PMLR article page. Venue comparison: the cited "ICML" corresponds to the official Proceedings of the 39th International…
Minor error	A Statistical Perspective on Retrieval-Based Models	2033			The PMLR proceedings page directly confirms the title, all three authors, the ICML proceedings venue, and publication year 2023. The ICML 2023 poster page independently confirms…
Minor error	Proximal Gradient Descent-Ascent: Variable Convergence under KŁ Geometry	2030	10.48550/arXiv.2102.04653	[illegible]	The ICLR poster record identifies the same title and all four authors as an ICLR 2021 paper. The arXiv record independently gives the same title and authors and says it…

Chimeric title.

Table 10: Chimeric title. No verified entries. 5 of 15 minor errors.

Verus	Title	Year	DOI	Matched DOI	Note
Minor error	Causal Inference Methods for Treatment Effect Estimation	2026	10.48550/arXiv.2602.18132	10.48550/arXiv.2602.18132	Preprint work at arxiv_static (authors match), but the cited title differs from the current official title (similarity 0.27) — typically a version retitle;…
Minor error	Multi-Modal Fusion for Visual Question Answering	2026	10.48550/arXiv.2602.18196	10.48550/arXiv.2602.18196	Preprint work at arxiv_static (authors match), but the cited title differs from the current official title (similarity 0.36) — typically a version retitle;…
Minor error	Federated Learning with Heterogeneous Client Distributions	2026	10.48550/arXiv.2602.18146	10.48550/arXiv.2602.18146	Preprint work at arxiv_static (authors match), but the cited title differs from the current official title (similarity 0.22) — typically a version retitle;…
Minor error	Neural Scene Representation and Rendering	2022	10.48550/arXiv.2602.18249	10.48550/arXiv.2602.18249	Preprint work at arxiv_static (authors match), but the cited title differs from the current official title (similarity 0.26) — typically a version retitle;…

Verus	Title	Year	DOI	Matched DOI	Note
Minor error	Retrieval-Augmented Generation for Long-Form Question Answering	2021	10.48550/arXiv.2602.12204	10.48550/arXiv.2602.12204	Preprint work at arxiv_static (authors match), but the cited title differs from the current official title (similarity 0.37) — typically a version retitle;…

Wrong venue.

Table 11: Wrong venue. 5 of 22 verified entries. 5 of 24 minor errors.

Verus	Title	Year	DOI	Matched DOI	Note
Verified	Dense Label Encoding for Boundary Discontinuity Free Rotation Detection	2021	10.1109/CVPR46437.2021.01553	10.1109/cvpr46437.2021.01553	CrossRef DOI lookup (score: 1.00)
Verified	Learning Local-Global Contextual Adaptation for Multi-Person Pose Estimation	2022	10.1109/CVPR52688.2022.01272	10.1109/cvpr52688.2022.01272	CrossRef DOI lookup (score: 1.00)
Verified	Learning Cross-Modal Retrieval With Noisy Labels	2021	10.1109/CVPR46437.2021.01536	10.1109/cvpr46437.2021.01536	CrossRef DOI lookup (score: 1.00)
Verified	MiniGPT-4: Enhancing Vision-Language Understanding with Advanced Large Language Models	2023	10.48550/arXiv.2304.10592	10.48550/arxiv.2304.10592	Matched arXiv id 2304.10592.
Verified	PET-NeuS: Positional Encoding Tri-Planes for Neural Surfaces	2023	10.1109/CVPR52729.2023.01212	10.1109/cvpr52729.2023.01212	CrossRef DOI lookup (score: 1.00)
Minor error	MonarchRT: Efficient Attention for Real-Time Video Generation	2026	10.48550/arXiv.2602.12271	10.48550/arxiv.2602.12271	Matched arXiv id 2602.12271. Cited DOI (10.48550/arXiv.2602.12271v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…
Minor error	Tiny Recursive Reasoning with Mamba-2 Attention Hybrid	2026	10.48550/arXiv.2602.12078	10.48550/arxiv.2602.12078	Matched arXiv id 2602.12078. Cited DOI (10.48550/arXiv.2602.12078v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…
Minor error	Meta-Sel: Efficient Demonstration Selection for In-Context Learning via Supervised Meta-Learning	2026	10.48550/arXiv.2602.12123	10.48550/arxiv.2602.12123	Matched arXiv id 2602.12123. Cited DOI (10.48550/arXiv.2602.12123v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…

Verus	Cite Title	Year	DOI	Matched DOI	Note
Minor error	Neutral Prompts, Non-Neutral People: Quantifying Gender and Skin-Tone Bias in Gemini Flash 2.5 Image and GPT Image 1.5	2022	10.48550/arXiv.2602.12133	10.48550/arXiv.2602.12133v1	arXiv id 2602.12133. Cited DOI (10.48550/arXiv.2602.12133v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…
Minor error	f-Domain Adversarial Learning: Theory and Algorithms	2021			The cited title and all four authors are confirmed by the PMLR proceedings record, which identifies the work as appearing in the Proceedings of the 38th International Conference…

Swapped authors.

Table 12: Swapped authors. 5 of 7 verified entries. 5 of 10 minor errors.

Verus	Cite Title	Year	DOI	Matched DOI	Note
Verified	Learning from Protein Structure with Geometric Vector Perceptrons	2021	10.48550/arXiv.2009.01411	10.48550/arXiv.2009.01411	arXiv id 2009.01411.
Verified	Calibrate Before Use: Improving Few-Shot Performance of Language Models	2021	10.48550/arXiv.2102.09690	10.48550/arXiv.2102.09690	arXiv id 2102.09690.
Verified	SDEdit: Guided Image Synthesis and Editing with Stochastic Differential Equations	2023	10.48550/arXiv.2108.01073	10.48550/arXiv.2108.01073	arXiv id 2108.01073.
Verified	The Mythos of Model Interpretability	2021	10.48550/arXiv.1606.03490	10.48550/arXiv.1606.03490	arXiv id 1606.03490.
Verified	Big Bird: Transformers for Longer Sequences	2023	10.48550/arXiv.2007.14062	10.48550/arXiv.2007.14062	arXiv id 2007.14062.
Minor error	Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback	2023	10.48550/arXiv.2307.15217	10.48550/arXiv.2307.15217	static matches title/DOI (score: 1.00), but cited author identities disagree with the official list. Unmatched cited author(s): Yoni Halpern. Official authors:…
Minor error	Understanding the Role of Training Regimes in Continual Learning	2023	10.48550/arXiv.2006.06958	10.48550/arXiv.2006.06958	static matches title/DOI (score: 1.00), but cited author identities disagree with the official list. Unmatched cited author(s): Dilan Gorur. Official authors:…
Minor error	Scaling Data-Constrained Language Models	2022	10.48550/arXiv.2305.16264	10.48550/arXiv.2305.16264	preprint work at arxiv_static (authors match), but the cited title differs from the current official title (similarity 0.52) — typically a version retitle;…

Verus	Title	Year	DOI	Matched DOI	Note
Minor error	Reinforcement Learning: An Introduction	2022			The MIT Press publisher page identifies the book as Reinforcement Learning: An Introduction, second edition, by Richard S. Sutton and Andrew G. Barto, published by The MIT Press…
Minor error	Align Your Latents: High-Resolution Video Synthesis with Latent Diffusion Models	2023	10.1109/CVPR52729.2023.02161	10.1109/CVPR52729.2023.02161	Open Access proceedings page confirms the cited work and provides its title, authors, CVPR venue, year, and pages 22563–22575. The official IEEE Xplore page also…

Preprint as published.

Table 13: Preprint as published. 5 of 24 verified entries. 5 of 8 minor errors.

Verus	Title	Year	DOI	Matched DOI	Note
Verified	Robust Audio-Visual Instance Discrimination	2021	10.1109/CVPR46437.2021.01274	10.1109/CVPR46437.2021.01274	CrossRef DOI lookup (score: 1.00)
Verified	Learning To Aggregate and Personalize 3D Face From In-the-Wild Photo Collection	2021	10.1109/CVPR46437.2021.01396	10.1109/CVPR46437.2021.01396	CrossRef DOI lookup (score: 1.00)
Verified	LLaMA: Open and Efficient Foundation Language Models	2023	10.48550/arXiv.2302.13971	10.48550/arXiv.2302.13971	arXiv id 2302.13971.
Verified	Constitutional AI: Harmlessness from AI Feedback	2023	10.48550/arXiv.2212.08073	10.48550/arXiv.2212.08073	arXiv id 2212.08073.
Verified	Robust Reinforcement Learning: A Case Study in Linear Quadratic Regulation	2021	10.1609/aaai.v35i10.17122	10.1609/aaai.v35i10.17122	Automatic CrossRef DOI lookup (score: 1.00)
Minor error	DEPTS: Deep Expansion Learning for Periodic Time Series Forecasting	2022			Title comparison: exact match to the official ICLR 2022 spotlight record. Venue comparison: the raw citation says “KDD,” but the direct official ICLR record and Microsoft…
Minor error	GREAD: Graph Neural Reaction-Diffusion Networks	2023			Title comparison: exact match. Venue comparison: the citation says IJCAI, but the direct PMLR proceedings article identifies the venue as the Proceedings of the 40th…
Minor error	Conjugate Energy-Based Models	2021			The official PMLR proceedings page directly confirms the cited title, all five authors, and year 2021. It identifies the venue as the Proceedings of the 38th International…

Verus	Title	Year	DOI	Matched DOI	Note
Minor error	Disambiguating Symbolic Expressions in Informal Documents	2021			Title comparison: exact match. Venue comparison: the raw citation says “WWW,” but the official OpenReview record identifies this paper as an ICLR 2021 Poster; arXiv’s record also…
Minor error	Adversarial Policy Learning in Two-player Competitive Games	2021			The PMLR proceedings page directly confirms the cited title, authors, and 2021 publication. Title comparison: exact match. Venue comparison: the citation says “ACL,” but the…

Hybrid fabrication.

Table 14: Hybrid fabrication. No verified entries. 5 of 7 minor errors.

Verus	Title	Year	DOI	Matched DOI	Note
Minor error	IBRNet: Learning Multi-View Image-Based Rendering	2021	10.1109/CVPR46437.2021.00466	10.1109/CVPR46437.2021.00466	The IEEE Computer Society Digital Library record (https://www.computer.org/csdl/proceedings-article/cvpr/2021/450900e688/1yeIC9n861q), inspected via URL context, confirms the…
Minor error	TensoRF: Tensorial Radiance Fields	2022	10.1007/978-3-031-19830-4_20	10.1007/978-3-031-19824-3_20	Title comparison: exact match. Venue comparison: Springer identifies it as a conference paper in Computer Vision – ECCV 2022, matching the cited ECCV venue. Author comparison:…
Minor error	ViViT: A Video Vision Transformer	2021	10.1109/ICCV48922.2021.00676	10.1109/ICCV48922.2021.00676	Title comparison: exact match. Venue comparison: the IEEE Computer Society record identifies the work as part of the 2021 IEEE/CVF International Conference on Computer Vision…
Minor error	Masked Autoencoders Are Scalable Vision Learners	2022	10.1109/CVPR52688.2022.01553	10.1109/CVPR52688.2022.01553	Title comparison: exact match. Venue comparison: CVPR 2022 is confirmed by the official CVPR Open Access record and IEEE Xplore. Author comparison: all six cited authors match…
Minor error	Mip-NeRF: A Multiscale Representation for Anti-Aliasing Neural Radiance Fields	2021	10.1109/ICCV48922.2021.00580	10.1109/ICCV48922.2021.00580	Confirmed by its IEEE Xplore publisher page, which gives the exact title, full author list, conference venue, year, and correct DOI…

Partial author list.

Table 15: Partial author list. 5 of 40 verified entries. 5 of 61 minor errors.

Verus	Title	Year	DOI	Matched DOI	Note
Verified	Role of Optimizer on Network Fine-tuning for Adversarial Robustness (Student Abstract)	2021	10.1609/AA[illegible]	10.1609/aaai.v35i17[illegible]	Static CrossRef DOI lookup (score: 1.00)
Verified	Contrastive Test-Time Adaptation	2022	10.1109/CV[illegible]	10.1109/CVPR52688.2022.00039	CrossRef DOI lookup (score: 0.99)
Verified	Matching Is Not Enough: A Two-Stage Framework for Category-Agnostic Pose Estimation	2023	10.1109/CV[illegible]	10.1109/CVPR52729.2023.01706	CrossRef DOI lookup (score: 0.99)
Verified	Contrastive Identity-Aware Learning for Multi-Agent Value Decomposition	2023	10.1609/AA[illegible]	10.1609/aaai.v37i10.26370	Static CrossRef DOI lookup (score: 0.99)
Verified	Retrieval-based Controllable Molecule Generation	2023	10.48550/AR[illegible]	10.48550/arXiv.2208.11126	arXiv id 2208.11126.
Minor error	Online Enhanced Semantic Hashing: Towards Effective and Efficient Retrieval for Streaming Multi-Modal Data	2022	10.1609/AA[illegible]	10.1609/aaai.v36i4.20346	DOI: the DOI landing page at https://doi.org/10.1609/aaai.v36i4.20346 is the AAAI publisher record. It confirms the exact title, all six authors, the venue, 2022…
Minor error	MoDem: Accelerating Visual Model-Based Reinforcement Learning with Demonstrations	2023	10.48550/arXiv.2212.05698	arXiv:2212.05698	Venue comparison: exact match. Venue comparison: the cited entry says ICLR, but direct records identify the work as arXiv:2212.05698 (submitted 12 Dec 2022); the Meta publication…
Minor error	Online Unrelated Machine Load Balancing with Predictions Revisited	2021			Direct primary evidence confirms the cited work: the ICML 2021 official poster record identifies the exact title and authors Shi Li and Jiayi Xian, and the PMLR proceedings PDF…
Minor error	SCoRe: Pre-Training for Context Representation in Conversational Semantic Parsing	2021			The direct ICLR 2021 poster record confirms the cited title, venue, year, and five-author list. The citation identifies the same work, but its author list omits three official…
Minor error	Knowledge distillation via softmax regression representation learning	2021			The official OpenReview record confirms the cited title and ICLR 2021 venue and identifies the full author list. The concrete error is that the citation omits two authors from…

Near miss title.

Table 16: Near miss title. 5 of 26 verified entries. 5 of 48 minor errors.

Verus	Title	Year	DOI	Matched DOI	Note
Verified	Online Enhanced Semantic Hashing: Towards Effective and Efficient Retrieval towards Streaming Multi-Modal Data	2022	10.1609/AAAI.V36I4.20346	10.1609/aaai.v36i4.20346	Static CrossRef DOI lookup (score: 0.98)
Verified	Training To Aggregate and Personalize 3D Face From In-the-Wild Photo Collection	2021	10.1109/CVPR46437.2021.01399	10.1109/cvpr46437.2021.01399	CrossRef DOI lookup (score: 0.98)
Verified	TransZero: Attribute-Guided Transformer towards Zero-Shot Learning	2022	10.1609/AAAI.V36I3.19999	10.1609/aaai.v36i3.19999	Static CrossRef DOI lookup (score: 0.97)
Verified	PAC-Bayesian Guarantees on Rate-Efficient Classifiers	2022			The PMLR publisher page directly confirms the cited title, authors, and 2022 publication. The raw citation gives the venue as “ICML,” which is the conference abbreviation; the…
Verified	Structured State Space Models for In Context Reinforcement Learning	2023		10.52202/075280-2036	Title comparison: the official record’s “In-Context” differs from the citation’s “In Context” only by hyphenation, a cosmetic variation. Venue comparison: “NeurIPS” is a short…
Minor error	UniT: Unified Multimodal Chain of-Thought Test-time Scaling	2026	10.48550/arXiv.2602.12279	10.48550/arXiv.2602.12279	arXiv id 2602.12279. Cited DOI (10.48550/arXiv.2602.12279v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…
Minor error	Self-supervised Learned Factor Analysis for Disentangling Utterance-level Speech Representations	2023	10.48550/arXiv.2305.08809	10.48550/arXiv.2305.08809	The arXiv record identifies the work as accepted by ICML 2023 and lists the four cited authors. The citation’s title says “Learned Factor Analysis,” whereas the official…
Minor error	Local Tuning Achieves Global Optimality in Multi-Agent Reinforcement Learning	2023			The direct PMLR proceedings record confirms the cited paper’s identity, authors, and year. The raw citation has a substantive title-word error: it says “Local Tuning” where the…
Minor error	Marginalized Stochastic Natural Gradients for Black-Box Variational Reasoning	2021			The official Proceedings of Machine Learning Research page identifies the same 2021 ICML work and lists the cited authors. However, the cited title ends in “Variational…
Minor error	Community Concealment from Self-supervised Graph Learning-Based Clustering	2026	10.48550/arXiv.2602.12250	10.48550/arXiv.2602.12250	Title comparison: the cited title says “Community Concealment from Self-supervised Graph Learning-Based Clustering,” while the arXiv primary record gives “Community Concealment…

Plausible fabrication.

Table 17: Plausible fabrication. No verified entries. All 2 minor errors.

Verus	Cite	Title	Year	DOI	Matched DOI	Note
Minor error		Optimization and generalization of gradient descent for shallow relu networks with minimal width	2021			The JMLR article page directly confirms the cited title and all four authors, but gives the publication as Journal of Machine Learning Research, 27(34):1—35, 2026. The cited…
Minor error		Language models of protein sequences at the scale of evolution enable accurate structure prediction	2023	10.1038/s41467-023-37674-3	10.1101/2022.07.20.500902	Title comparison: the cited title exactly matches the bioRxiv record. Venue comparison: the raw citation identifies Nature Communications and gives DOI…

Arxiv version mismatch.

Table 18: Arxiv version mismatch. 5 of 45 verified entries. 5 of 25 minor errors.

Verus	Cite	Title	Year	DOI	Matched DOI	Note
Verified		SCAFFOLD: Stochastic Controlled Averaging for Federated Learning	2023	10.48550/arXiv.1910.06378	10.48550/arXiv.1910.06378	arXiv id 1910.06378.
Verified		Instance-Aware Domain Generalization for Face Anti-Spoofing	2021	10.1109/CVPR46437.2021.00501	10.1109/CVPR52729.2023.01579	CrossRef DOI lookup (score: 0.95)
Verified		ACSeg: Adaptive Conceptualization for Unsupervised Semantic Segmentation	2022	10.1109/CVPR46437.2021.00692	10.1109/CVPR52729.2023.00692	CrossRef DOI lookup (score: 1.00)
Verified		PET-NeuS: Positional Encoding Tri-Planes for Neural Surfaces	2021	10.1109/CVPR46437.2021.01212	10.1109/CVPR52729.2023.01212	CrossRef DOI lookup (score: 0.95)
Verified		Matching Networks for One Shot Learning	2023	10.48550/arXiv.1606.04080	10.48550/arXiv.1606.04080	arXiv id 1606.04080.
Minor error		Prototypical Networks for Few-shot Learning	2023	10.5555/3295222.3295383		Official NeurIPS proceedings BibTeX record identifies this title and author list as a 2017 work (volume 30). The official NeurIPS paper PDF also identifies it as part of the…
Minor error		Structured State Space Models for In-Context Reinforcement Learning	2022		10.52202/075280-2036	Official NeurIPS proceedings page and DOI landing page confirm this title and the cited authors. The proceedings record identifies it as a NeurIPS 2023 paper in Advances in…
Minor error		A technical curriculum on language-oriented artificial intelligence in translation and specialised communication	2023	10.48550/arXiv.2602.12251	10.48550/arXiv.2602.12251	arXiv id 2602.12251. Cited DOI (10.48550/arXiv.2602.12251v1) does not resolve (CrossRef/DataCite/doi.org) — cannot remain verified (bad/unregistered DOI). Identity…

VerusCite	Title	Year	DOI	Matched DOI	Note
Minor error	Hypergraph Propagation and Community Selection for Objects Retrieval	2022			The official NeurIPS proceedings page confirms the cited title, authors, and NeurIPS venue, and identifies the work as part of NeurIPS 2021. The citation gives 2022, a concrete…
Minor error	Rigid Body Flows for Sampling Molecular Crystal Structures	2022			Title comparison: exact match to the official PMLR proceedings page. Venue comparison: the citation says NeurIPS, but the publisher/proceedings record identifies the work as part…

References

Kim, Kayoon, Wieland Brendel, and Patrik Reizinger. 2026. "HALLMARK: How to Diagnose LLM Citation Verifiers and When to Deploy Them." https://arxiv.org/abs/2607.18360.

Reizinger, Patrik, Kayoon Kim, and Wieland Brendel. 2026. "HALLMARK: Citation Hallucination Detection Benchmark." https://github.com/rpatrik96/hallmark.

Wheeler, Andrew P. 2026a. "VerusCite." https://veruscite.com.

———. 2026b. "VerusCite API." https://veruscite.com/docs/api/.